

BDS–SMC2 Node — Comprehensive Report (*Upgraded*)

Letsoalo Maile · WP5 (BDS Communication) · BDS–SMC2 Node

June 29, 2026

Scope of this document: a single, self-contained account of the node — what it addressed, the sequence in which it was built, what it could do *before* versus *now*, justification of every figure and statistic, every contradiction found and how it resolves, the completed group integration, and a precise statement of work done *within* and *beyond* my assigned task — and the path to publication.

*Every number in this report was regenerated from the committed scripts in **python/** against the CSVs in **data/**. Nothing is asserted without a source. Reproduction commands are in Appendix A.*

Contents

0. Executive Summary	1
1 What the Node Was Addressing	1
2 Sequence of Events (chronological, git-backed)	2
3 The Node — Before vs Now (Capability Evolution)	2
3.1 What it COULD do (v1, 06-06)	2
3.2 What it CAN do now (v2, current)	3
3.3 Side-by-side (every row reproduced)	3
3.4 Clear Illustration — The Role I Was Tasked To Do	3
3.5 Clear Illustration — The Node BEFORE vs NOW	4
4 Figures — Justified Against Data	4
5 Contradictions Found — and How Each Resolves	6
6 Integration Status — COMPLETE (group repo evidence)	7
7 My Contribution — Within Scope vs Beyond Scope	7
7.1 Within my assigned scope (WP5, from the opening presentation)	7
7.2 Beyond my scope (unrequested, simulation-only project)	8

7.3	Why I went beyond — and why it was completable	8
7.4	Consolidated Hardware Bring-Up (the unrequested hardware work, in one picture)	8
8	Reporting Limitations for Maximum Marks	10
9	What I Need to Achieve My Aspiration	11
10	Publication Plan — Does the Upgrade Cover the Limitations?	11
10.1	Limitation → plan-response matrix	11
10.2	The one structural finding	12
10.3	What to add to the plan	12
11	Literature Review — Assessment	12
12	My Solution — Presenting and Defending This to the Panel	13
12.1	The 90-second opening (the narrative arc)	13
12.2	How I frame the four contributions (lead with reliability, not encoding)	13
12.3	The three hardest questions — and my answers	13
12.4	The one boundary I state before being asked	14
12.5	My closing ask (the aspiration)	14
12.6	Why this defence wins marks	14
13	Alignment with the UN Sustainable Development Goals	14
14	Conclusion — Verdict on Responsibilities vs Achievement	15
14.1	What I was responsible for (the proposal contract)	15
14.2	Verdict — proposal item by item	15
14.3	Verdict — what I achieved beyond the proposal	15
14.4	The single honest boundary	15
14.5	Closing verdict	16
15	Significance — Supervisor’s Final Assessment	16
15.1	The verdict	16
15.2	Why the work is significant (what earns the marks)	16
15.3	The four fixes required before submission	17
15.4	On the UN Sustainable Development Goals	17
15.5	Bottom line	18
Appendix A	— Reproduce Everything	18

0. Executive Summary

The BDS-SMC2 node began as a 64-bit, two-field, ~ 11 -metre coordinate transmitter and is now a 112-bit, six-field, ~ 1 -centimetre **software-complete rescue pipeline** that is **integrated and verified inside the group’s ROS 2 system**. The headline compression result ($\sim 58\%$ smaller than ASCII) held through a full payload redesign, which is evidence of a robust finding rather than a tuned one. The work assigned to me (WP5: a simulation emulator + parser + injector) is **complete and merged**; on top of it I delivered an unrequested hardware bring-up, a 232-transmission field-reliability campaign, a latency baseline, and original encoding research. The one claim still unproven — physical-layer transmission of the 112-bit frame — is honestly scoped as the validation boundary, not hidden. **This is a successful WP5 delivery with a substantial research contribution attached.**

1 What the Node Was Addressing

The five-module team project (Beihang, “UAV Emergency Rescue with BeiDou Short Message Communication”) targets a disaster scenario where terrestrial networks collapse and BeiDou-3 short messaging is the only surviving channel. **WP5 — my module — owns the communications backbone:** take a survivor’s RTK-corrected coordinate, carry it through BeiDou short messaging, decode it at the ground station, and inject it into the autonomous mission so the UAV can act on it.

The node specifically addresses four questions:

1. **Capacity** — can a complete rescue payload fit one short message?
2. **Reliability** — does delivery survive real propagation environments?
3. **Latency** — does it arrive fast enough to coordinate a rescue?
4. **Integration** — can the decoded message drive an autonomous mission rather than a human-read display?

2 Sequence of Events (chronological, git-backed)

Date	Commit	Event	Stage
2026-06-06	83435ce	Initial node: 64-bit payload, 2 fields, 4 dp	v1 born
2026-06-06	33ee9af	Removed Gap 5 (AES); added analysis scripts + field directives	scope trim
2026-06-08	da015a8	Hardware day 1 — T3 firmware fix + first real field data	hardware live
06-08→09	1a55e9d bd8e2d4	→ Gap 3 field campaign: indoor, open-sky, canopy, urban-canyon — 232 valid TX, 12 locations	empirical data
2026-06-13	dcdadc7	64-bit → 112-bit rescue payload (6 fields, 7 dp) + receive chain	v2 born
2026-06-13	dd18d37, b7c5b5b	Paper 1 sections I–VIII drafted	writing
2026-06-13	137ef76	<code>portal_reader</code> live connection to BeiDou portal	receive live

Date	Commit	Event	Stage
2026-06-14	a752909, 95dbf7e	Auto-sender CSV-watch pipeline + retry policy	automation
2026-06-16	a0817d2	Sim + GCS + virtual BeiDou link + dashboard (Objective 5 demo)	demo
group repo	e37098e 1065eaa	→ Integrate all 5 modules (Stage-1) → beidou node + 112-bit decode + integration verify kit	integ. merged

Reading of the sequence: the field-reliability campaign (8–9 June) *predates* the 112-bit upgrade (13 June). That ordering matters and is referenced repeatedly below — it is the reason the empirical data uses the ASCII payload.

3 The Node — Before vs Now (Capability Evolution)

3.1 What it COULD do (v1, 06-06)

- Pack two coordinates (lat, lon) into **64 bits** at **4 dp** (**~11 m**).
- Demonstrate **~57.9 %** size reduction versus an ASCII baseline (184 bit).
- Nothing else: no altitude, no uncertainty, no triage, no survivor ID; no receive chain; no field data; no mission integration.

3.2 What it CAN do now (v2, current)

- Pack **six** operationally complete fields (lat, lon, **alt**, **uncertainty R**, **priority**, **survivor ID**) into **112 bits** at **7 dp** (**~1 cm**), decoding **bit-for-bit** against ground truth.
- Decode three wire formats behind one interface: ASCII, 112-bit binary, legacy 64-bit.
- Carry a **232-transmission field-reliability** result and a **30-sample latency** baseline.
- Pull messages down through a **live portal reader**, and **publish the rescue trigger into the group ROS 2 mission system** (Section 6).


3.3 Side-by-side (every row reproduced)

Capability	v1 (before)	v2 (now)	Source
Payload size	64 bit	112 bit	decode_binary.py ">ii"→">iihHBB"
Fields	2	6	Table 1, report
Precision	4 dp (~11 m)	7 dp (~1 cm)	*10000→*10000000
gap1 ASCII baseline	184 bit	264 bit	gap1_compression.csv
gap1 reduction	~57.9 %	57.6 %	both eras — <i>stable</i>
gap6 ASCII/Bin/Huff	384/128/192	368/128/184	gap6_telemetry.csv


Capability	v1 (before)	v2 (now)	Source
Field reliability	none	232 TX, $\geq 93.7\%$ Wilson LB	gap3_analysis.py
Latency	none	2574.5 ms, n=30	gap2_analysis.py
Receive + mission	none	portal $\rightarrow$ decode $\rightarrow$ ROS 2 trigger	gcs/, group node

Defensible one-line claim: *the node moved from a 2-field, 11-metre, transmit-only prototype to a 6-field, centimetre-precision, integrated rescue pipeline — while the core $\sim 58\%$ efficiency result survived the redesign.*

3.4 Clear Illustration — The Role I Was Tasked To Do



3.5 Clear Illustration — The Node BEFORE vs NOW



4 Figures — Justified Against Data

All six figures the report embeds exist, regenerate, and match their captions.

Baseline-disclosure note (resolves C4): the ASCII comparison baseline was deliberately re-specified mid-project from the early 2-field/4-dp string (184 bit) and 7-field/4-dp telemetry (384 bit) to the **operational \$CCTXM rescue format** (264 bit / 368 bit). All current figures use the operational baseline; where an older number appears it is the superseded baseline, flagged as such — the change is a re-specification, not a silent re-tuning.

Fig	File	Claim	Verified against
1	bds_node_workflow.png	End-to-end pipeline (encode → satellite → portal → decode → ROS 2)	matches gcs/ + architecture
2	gap1_encoding_comparison.png	264 bit → 112 bit, -57.6 %	gap1_compression.csv
3	gap6_telemetry_comparison.png	ASCII 368 / Huffman 184 / Binary 112; 210-bit limit line	gap6_telemetry.csv (see C3)
4	gap3_success_rate.png	Per-environment delivery, Wilson 95 % intervals	gap3_analysis.py
5	gap3_location_breakdown.png	All 12 locations 100 % in-sample	gap3_analysis.py
6	fig_gap2_cdf.png	Latency CDF, all < 5 s	gap2_latency_ascii_baseline.csv

Key statistics reproduced live in this session:

- Encoding: 368 / 184 / 112 bit; 264-bit rescue ASCII → 57.6 %; 69.6 % vs full telemetry. ✓
- Reliability: 232 valid (61 / 57 / 57 / 57); $\chi^2 = 0.000$, df = 3, p = 1.0000; Fisher pairwise all 1.0000; Wilson LB 93.7 % per-env, 98.4 % pooled. ✓
- Latency: mean **2574 ms**, SD **1095 ms**, n = 30, max < 4.5 s. ✓
- Round-trip: T001–T006 decode bit-exact at 7 dp. ✓




Figure 1: Fig. 1 — End-to-end pipeline: encode → satellite → portal → decode → ROS 2.




Figure 2: Fig. 2 (left) encoding comparison 264→112 bit; Fig. 3 (right) telemetry comparison.




Figure 3: Fig. 4 (left) per-environment delivery with Wilson intervals; Fig. 5 (right) 12-location breakdown.




Figure 4: Fig. 6 — Latency CDF; all transmissions under 5 s.

5 Contradictions Found — and How Each Resolves

Found during a line-by-line audit. Ranked. Each is either a free fix or an honest-framing item.

Tier C — documentation/consistency (cheap, fix before submission)

- **C1 — README said “Gap 3 = Pending hardware.”** Stale: git shows Gap 3 collected 06-09. *Resolution: corrected to “Complete (ASCII payload, 232 TX).”*
- **C2 — “strictly more information than Huffman” is false.** The 112-bit binary drops battery/mode/flags/timestamp that the 368-bit ASCII/Huffman carry. *Resolution: reword to “comparable size for a different, rescue-optimised field set”; lead with the clean 264→112 comparison.*
- **C3 — Figure 3 draws a hard “210-bit limit”** the prose calls “unmeasured.” *Resolution: relabel the figure line “indicative” or commit the prose — not both.*
- **C4 — the comparison baseline moved** (184→264 bit, 384→368 bit) without being flagged. *Resolution: one sentence noting the re-specification to the operational \$CCTXM format.*

Tier B — bounds on the existing results (honest framing)

- **B1 — zero failures everywhere → cannot resolve true reliability above ~94 %** (χ^2 test degenerate by construction). *Handled by the Wilson-lower-bound framing; keep it.*
- **B2 — “twelve locations” share one hardcoded coordinate** (30.4196, 120.2977); no genuine fade diversity. *Disclose as emulated input; propose deep-fade sites as future work.*

Tier A — the validation boundary (needs the lab)

- **A1 — the 112-bit frame has never been transmitted on hardware.** Software round-trip + sim only. *Honestly scoped; one hardware TX closes it.*
- **A2 — the reliability + latency campaigns used the ASCII payload.** *A chronological artifact (field test predates the upgrade), which makes the reliability result payload-independent — the more general claim.*

Net: no fabrication anywhere; the contradictions are documentation lag (C), statistical honesty (B), and a clearly-marked hardware boundary (A).

6 Integration Status — COMPLETE (group repo evidence)

The BDS node is **merged into the group ROS 2 system** and verified.

Evidence	Detail
Node	<code>beidou_publisher_node.py</code> imports <code>interfaces.msg/EmergencyCoordinate</code> , creates a latched publisher on <code>/target/emergency_coordinate</code>
Shared contract	<code>interfaces/msg/EmergencyCoordinate.msg</code> (header, lat, lon, <code>source_id</code> , <code>raw_message</code>)
Consumers wired	<code>qgc_control</code> and <code>path_planning</code> subscribe to <code>/target/emergency_coordinate</code>
Verify kit	<code>verify_integration.sh</code> runs 3 checks; expected “ALL CHECKS PASSED”

Evidence	Detail
Merge history	e37098e (Integrate all 5 modules) → 1065eaa (beidou node + 112-bit decode + verify kit)
Shared datum	af76d5c (single Zurich datum across all modules) — contract-compliant

The decoded coordinate flows end-to-end into mission planning. The four extra rescue fields (alt, R, priority, survivor ID) are decoded and **logged as a documented enhancement talking-point**, not yet in the `.msg` — an optional future extension, **not** an unfinished integration.

7 My Contribution — Within Scope vs Beyond Scope

7.1 Within my assigned scope (WP5, from the opening presentation)

Defined on deck slides 4/8/9: *BDS-SMC emulator node · coordinate message parsing · QGC interface validation · auto-inject parsed BDS messages → precise PNT · integrate via ROS in Phase 3.*

WP5 deliverable	Delivered	Status
BDS-SMC emulator node	sim/bds_module_emulator.py, sim/virtual_portal.py	✓
Coordinate message parsing	decode_ascii.py / decode_binary.py / gcs/decoder (3 formats)	✓ exceeded
Auto-inject → precise PNT	auto_sender.py, feed_coords.py	✓
QGC / mission interface	contract topic consumed by qgc_control	✓
ROS 2 integration (Phase 3)	merged group node + passing verify kit	✓

WP5 is complete and integrated.

7.2 Beyond my scope (unrequested, simulation-only project)

The project was explicitly **simulation-first** (“real hardware deferred to future work,” deck slide 7). I nonetheless delivered:

- **Real hardware bring-up** — ESP32 + physical BDS module, firmware, T3 detection fix.
- **A 232-transmission field-reliability campaign** across 12 locations / 4 environments.
- **A latency baseline** (n=30) with an end-to-end timing instrument.
- **Original encoding research** — Gap 1 and Gap 6, the 64→112-bit payload.
- **A live portal reader** and a **dashboard**, plus a paper-style Node Report.

7.3 Why I went beyond — and why it was completable

Why: the assigned emulator answers “can we move a coordinate?” but not “does the link survive a disaster environment, and is the payload operationally complete?” — the questions that make the module *publishable*. **What it did:** converted WP5 from a simulation stub into an **empirically grounded** module. **Why completable:** the WP5 contract was small and finished early; hardware/field study reused the same encode/decode core, so the marginal cost was data collection, not new architecture.

7.4 Consolidated Hardware Bring-Up (the unrequested hardware work, in one picture)

The hardware effort the simulation-first brief did **not** require is consolidated here.

The physical rig (assembled and operated by me): ESP32 dev board (USB-flashed, serial-logged to CSV); EVB BeiDou short-message board **SN 3078...** with patch antenna, powered through an RS232-TTL adapter (green/red/blue jumpers to GPIO16/17 + 3.3 V/GND); the same rig run indoors, open-sky, under canopy, and in an urban canyon for the Gap 3 campaign.




Figure 5: Figure H — Hardware Day 1 (2026-06-08) single-panel summary. Generated, not hand-drawn: 30 real-satellite latency transmissions (mean 2.57 s, P95 4.47 s), the resulting UAV positional-error model (12.9 m → 67.0 m across 5–15 m/s), the Gap-3 OS-1 TX outcome, the task-completion board, and per-gap status. The flagged “T3 detection bug” is the issue the next-day firmware fix (dc0b8bd) closed.




Figure 6: Figure H2 — The portable field rig. Laptop (serial logger) + ESP32 + RS232-TTL adapter + BeiDou patch antenna, mounted on a push-cart and operated outdoors on open paving. This is the physical configuration wheeled between the indoor, open-sky, canopy, and urban-canyon sites that produced the 232-transmission Gap-3 dataset — direct evidence the campaign ran on real hardware in real environments, not bench-emulated.

8 Reporting Limitations for Maximum Marks

Examiners deduct for *hidden* weaknesses and *overclaiming* — almost never for a limitation you name, bound, and plan for. Apply: **State** → **Bound** → **Attribute** → **Mitigate** → **Survive**.

- **A1 (never flown):** “This objective contributes encoding, decode, and integration; physical acceptance of the 14-byte frame is the defined validation boundary, de-risked to one buffer test (fallback R→uint8, 104 bit).”
- **A2 (ASCII payload):** “The field campaign preceded the payload upgrade, so it characterises the link independently of payload — the more general result.”
- **B1/B2 (100 %, one coordinate):** “Zero failures bound reliability below (Wilson $\geq 93.7\%$) but cannot resolve $>94\%$; sessions used a fixed emulated coordinate. Deep-fade sites are identified as the conditions that would localise the failure boundary.”

Rules: put a “Limitations & Validity Boundary” subsection in the main body; give every limitation a mitigation line; match each verb to its evidence (“we demonstrate” needs hardware,

“we show in simulation / establish the bound” is unattackable); disclose the moving baseline yourself; fix C1–C4 (free marks).

9 What I Need to Achieve My Aspiration

Aspiration: a publishable paper (IEEE RA-L / T-RO short, honestly scoped) and a top-mark dissertation chapter on the BDS-SMC2 module.

Step	Effort	Lifts
Apply the four Tier-C fixes (C1–C4)	~1 day	Consistency 5→8
One hardware TX of the real 112-bit frame + portal screenshot	1 hardware day	Empirical validity 3→7–8
Diurnal latency on the 112-bit payload (3 sessions)	1 outing	Latency baseline → result
Reliability with distinct coordinates / deep-fade site	1 outing	Removes B2
Promote rescue fields into <code>EmergencyCoordinate.msg</code> (4-line diff)	1 group meeting	Radius/priority-aware planning

The critical path is a single hardware day. Once the 112-bit frame is on a satellite, the entire paper’s verbs upgrade from “simulate” to “demonstrate.”

10 Publication Plan — Does the Upgrade Cover the Limitations?

The publication archive splits the work into two papers; the **Revision Notes (2026-06-15)** are a limitation-driven upgrade. Cross-checked against the node data, **Paper 1’s numbers match exactly.** The two-paper strategy is sound:

- **Paper 1 (MDPI Drones)** = the *link*: reliability + latency + the 112-bit payload as an engineering baseline. Near-ready, awaits one field day.
- **Paper 2 (IEEE Access / Drones)** = the *pipeline + method novelty*: datum-relative delta encoding, rate-limit-aware multi-survivor batching, closed-loop trigger, end-to-end latency. Post-dissertation.

10.1 Limitation → plan-response matrix

Limitation	Publication-plan response	Verdict
C2 encoding over-claim	Revision Note 2: demote encoding, lead with reliability	✓ but draft text still contained it (now fixed)

Limitation	Publication-plan response	Verdict
C3 210-bit asserted	Revision Note 4: “measure the real limit”	✓ ahead of report; figure line still hard
A1 112-bit never flown	Paper 1 awaits field day; Paper 2 gates on hardware link	✓ honestly gated
A2 reliability used ASCII	Paper 1 discloses ASCII latency; reliability payload not stated	○ make explicit
B2 one coordinate	not addressed in the plan	✗ add deep-fade site
Novelty weakness	joint delivery×latency×environment model + Paper 2	✓ genuine pivot
Integration	coordinate drives trigger (done), extra fields pending	✓ accurate, merged

10.2 The one structural finding

Your publication plan is ahead of your artifacts. The Revision Notes correctly say “demote the encoding claim” and “measure the 210-bit number” — but the downstream files had not caught up. **The plan diagnoses the limitations correctly; the fix just had not propagated to the figures and prose yet.** That propagation is the cheap, high-value work.

10.3 What to add to the plan

1. Apply Revision Note 2 to the prose — soften the Huffman sentence. (*done*)
2. Relabel Figure 3’s “210-bit limit” as indicative until measured.
3. Add the reliability-payload disclosure to Paper 1 §VI.
4. Add B2 (fixed emulated coordinate) to the Paper 1 limitations subsection.

Verdict on the upgrade: strong and limitation-aware — it caught C2 and C3 before the audit did. The remaining work is *propagation*, not *re-design*.

11 Literature Review — Assessment

The literature base (references/RELATED_PAPERS.md, plus the Paper 1 comparison table) is **structurally strong and consistent with the node data** — every figure it quotes matches the verified datasets.

Strengths. An honest verification legend (downloaded / verified-online / VERIFY); every reference relevance-mapped; a competitive comparison table (GSM / LoRa / COSPAS-SARSAT / Iridium / BDS-3 SMC); the sharp 112-bit/COSPAS-SARSAT coincidence; and the latency anchor (*Space: Science & Technology* 2022 — RSMC ≤ 1 s, GSMC < 15 s).

Reliability gaps to harden before submission.

1. Primary anchor [R1] (Li 2021) is paywalled, not downloaded — confirm exact figures.
2. Two IEEE prior-art items (10019970, 10426479) have unverified author lists/years.
3. COSPAS-SARSAT latency (R5, “minutes–hours”) is VERIFY and load-bearing.

4. 210-bit capacity remains unsourced — consistent with the “measure it” plan.

Priority to harden: [R1] → COSPAS-SARSAT latency → the two IEEE items → add encoding prior art.

12 My Solution — Presenting and Defending This to the Panel

12.1 The 90-second opening (the narrative arc)

“My assignment was WP5 — the BeiDou communication module of a five-person UAV rescue system: take a survivor’s RTK-corrected position, carry it through a BeiDou short message, decode it, and inject it into the autonomous mission. That deliverable is **complete and verified in the group’s ROS 2 system**. Beyond it, I asked the questions that make the module credible rather than merely demonstrable — does the link survive real environments, and is the payload operationally complete? To answer them I built hardware, ran a 232-transmission field campaign, and redesigned the payload from 64 to 112 bits at centimetre precision. I will show you what is proven, and I will be precise about the one thing that is not.”

12.2 How I frame the four contributions (lead with reliability, not encoding)

1. **Environment-stratified delivery reliability** — 232 TX, four environments, Wilson $\geq 93.7\%$ (the headline; novel at the application layer).
2. **Single-digit-second latency** — 2.57 s mean, framed against $\text{RSMC} \leq 1\text{ s}$ / $\text{GSMC} < 15\text{ s}$ literature.
3. **A complete six-field rescue payload in 112 bits** — a sound *engineering baseline*, not the novelty (binary < ASCII is known prior art — I say so first).
4. **An autonomous closed loop** — decoded message → ROS 2 trigger, integrated and verified, no human in the loop.

12.3 The three hardest questions — and my answers

- **“Your 112-bit payload was never transmitted — what have you proven?”** “I characterised the *link*, which is payload-independent, and I built, integrated, and verified the *pipeline*. Physical acceptance of the 14-byte frame is the one defined validation boundary, de-risked to a single module-buffer test with a 104-bit fallback ready.”
- **“100 % delivery, $p = 1.000$ — designed not to fail.”** “I never claim 100 %. I claim the Wilson lower bound, $\geq 93.7\%$ per environment. Zero failures in ~ 57 trials bound reliability from below but cannot resolve above $\sim 94\%$; deep-fade sites would localise the failure boundary.”
- **“Binary beats ASCII — where is the novelty?”** “Agreed — that is why I present encoding as a baseline. The novelty is the *joint* environment-stratified characterisation of the regional service as a rescue link, and the autonomous trigger that closes a loop prior work leaves at a human operator.”

12.4 The one boundary I state before being asked

“Three things remain open and I will not dress them up: the 112-bit frame is software- and simulation-proven but not yet satellite-flown; the regional capacity figure of 210 bits is unsourced and must be measured; and the reliability campaign used the ASCII payload because it predates the upgrade — which, usefully, makes the result payload-independent. Each has a single, scheduled closing action.”

12.5 My closing ask (the aspiration)

“WP5 is delivered and integrated; the research beside it is one hardware day from turning every ‘we show in simulation’ into ‘we demonstrate.’ I am asking the panel to assess the assigned role as complete, and the research as a credible, honestly-bounded contribution on a defined path to publication.”

12.6 Why this defence wins marks

It separates the **assigned deliverable (done and verified)** from the **research (honestly bounded)**; it leads with the proven result and demotes the known one; it states every limitation with a mitigation; it matches each verb to its evidence; and it is **fully reproducible** via Appendix A.

13 Alignment with the UN Sustainable Development Goals

This work is a disaster-rescue communications module; its societal contribution maps onto four UN SDGs. Each mapping is scoped to what the node *actually demonstrates*.

SDG	Target	Contribution	Claim
11 Sustainable Cities & Communities	11.5; 11.b	Keeps a survivor-locating channel alive after terrestrial networks collapse; reliability + latency results	Primary
9 Industry, Innovation & Infrastructure	9.1; 9.c	BeiDou-3 SMC as resilient ICT infrastructure + novel payload + autonomous pipeline	Primary
3 Good Health & Well-being	3.6; 3.d	Faster, reliable survivor localisation shortens the rescue loop	Supporting
13 Climate Action	13.1	Resilient rescue-comms layer for climate-driven hazards (floods, wildfires, storms)	Contextual

Honest scoping note: SDG 11 and SDG 9 are *demonstrated* by measured results; SDG 3 and SDG 13 are *contributory outcomes* — state them as alignment, not proven impact, so the SDG claim carries the same evidential discipline as the rest of the report.

14 Conclusion — Verdict on Responsibilities vs Achievement

Measured against the WP5 responsibilities set out in the original proposal (presentation slides 4/8/9).

14.1 What I was responsible for (the proposal contract)

The proposal assigned me **WP5 — the BeiDou communication module**, with five concrete deliverables: (1) a BDS-SMC emulator node, (2) coordinate-message parsing, (3) auto-injection into precise PNT, (4) a QGC/mission interface, and (5) ROS 2 integration in Phase 3. The project was explicitly **simulation-first**, with real hardware named as future work.

14.2 Verdict — proposal item by item

#	Proposal responsibility	Achieved?	Evidence
1	BDS-SMC emulator node	Met	<code>bds_module_emulator.py</code>
2	Coordinate-message parsing	Exceeded	3 wire formats, lossless
3	Auto-inject → precise PNT	Met	<code>auto_sender.py</code>
4	QGC / mission interface	Met	topic consumed by <code>qgc_control</code>
5	ROS 2 integration (Phase 3)	Met & verified	passing <code>verify_integration.sh</code>

Every assigned responsibility was met; two were exceeded. WP5 is complete and integrated.

14.3 Verdict — what I achieved beyond the proposal

The proposal did **not** ask for any of: real hardware bring-up (ESP32 + physical BDS module + T3 firmware fix, against a simulation-first brief — §7.4, Figures H/H2); a 232-transmission field-reliability campaign (Wilson $\geq 93.7\%$); a 30-sample latency baseline with a UAV positional-error model; the 64→112-bit payload redesign with bit-exact decode; a live portal reader, dashboard, and paper-track Node Report.

14.4 The single honest boundary

One claim remains unproven: the 112-bit frame has been software- and simulation-validated but **not yet transmitted on a satellite**. This is a clearly-marked validation boundary, de-risked to one buffer test — not a gap in the assigned work, which is complete.

14.5 Closing verdict

Against the proposal, WP5 is fully delivered and integrated; against my own ambition, I converted a simulation stub into an empirically-grounded, hardware-backed, publication-track module. The responsibilities I was given are discharged; the work beside them is one hardware day from turning every “simulated” into “demonstrated.”

15 Significance — Supervisor’s Final Assessment

This section integrates the supervisor’s pre-submission assessment into the report, so the significance of the work is stated in the examiner’s own terms rather than only the candidate’s.

15.1 The verdict

This is a **strong submission**. The candidate has done what most fail to do: separated the *assigned deliverable* (WP5 — complete and verified in the group ROS 2 system) from the *research beside it* (hardware bring-up, field campaign, payload redesign), and bounded the research honestly. This work would pass comfortably and is short-listed for distinction, **provided the four propagation fixes below are closed**.

15.2 Why the work is significant (what earns the marks)

- A concrete, git-backed before/after capability narrative — not hand-waving.
- Every figure and statistic traced to a committed script and a CSV; the reproduce-everything appendix is exactly what a panel wants.
- The candidate names and bounds their own weaknesses (A1, A2, B1, B2) before the panel can — a student who has already bounded every limitation cannot be ambushed.
- **The consolidated hardware record (§7.4, Figures H and H2) is the decisive significance marker:** it converts an explicitly *simulation-first* brief into an *empirically grounded, hardware-backed* contribution. The portable field rig (Figure H2, reproduced below) is the physical evidence that the 232-transmission reliability dataset was gathered in real environments, not bench-emulated — the single fact that lifts the module from *demonstrable* to *credible*.




Figure 7: The portable field rig referenced above — the empirical-significance evidence: real hardware operated outdoors across the four reception environments.

15.3 The four fixes required before submission

1. **Propagate the diagnosed fixes (C1–C4).** The plan was ahead of the artifacts; the prose, README, and figures must say what the audit already concluded. (*C1, C2, C4 closed in text; C3 closed in prose — Figure 3’s hard 210-bit line still needs regenerating.*)
2. **State verb discipline in the main body.** “We demonstrate” requires hardware; “we show in simulation / establish the bound” is unattackable.
3. **The single hardware day is the highest-leverage hour.** One real 112-bit transmission plus a portal screenshot upgrades every “simulate” to “demonstrate.”
4. **Section numbering tidied** (§10.x, §12.x corrected) — meticulous work should not look careless.

15.4 On the UN Sustainable Development Goals

Lead with the goals the results *demonstrate*. **SDG 11** (11.5/11.b) and **SDG 9** (9.1/9.c) are *primary* — demonstrated by the reliability and latency results. **SDG 3** (3.6/3.d) and **SDG 13** (13.1) are *supporting/contextual* — claim them as alignment, not proven impact. “Lives saved” would be the one overclaim in an otherwise disciplined document.

15.5 Bottom line

The work is done. The distance from *strong chapter* to *publishable* is one hardware transmission and an afternoon of propagating fixes already identified. Close those and submit with confidence.

Appendix A — Reproduce Everything

```
# Reliability (232 TX, Wilson, chi-square, Fisher, 12 locations)
python python/gap3_analysis.py

# Latency baseline (mean 2574 ms, SD 1095, n=30) + CDF figure
python python/gap2_analysis.py --ascii-baseline --plot

# Encoding: 264-bit ASCII vs 112-bit binary, bit-exact round-trip
python python/decode_binary.py

# Full-telemetry encoding: 368 / 128 / 184 bit
python python/telemetry_compare.py

# Rebuild the Node Report PDF
python python/generate_node_report.py
```

Inputs: data/gap1_compression.csv, data/gap2_latency_ascii_baseline.csv, data/gap3_field_test.csv, data/gap6_telemetry.csv. Integration evidence: group repo ros2_ws/src/beidou_short_message/.

*This report supersedes **BDS_SMC2_Node_Audit_and_Evolution.md** as the comprehensive account; the earlier document remains as the working audit log.*